

# CS 5220

## Parallelism and Locality in Simulations

---

David Bindel

2024-09-17

## Lumped Parameter Models


---

Examples include:


- SPICE-level circuit simulation
  - nodal voltages vs. voltage distributions
- Structural simulation
  - beam end displacements vs. continuum field
- Chemical concentrations in stirred tank reactor
  - concentrations in tank vs. spatially varying concentrations

- Typically ordinary differential equations (ODEs)
- Constraints: differential-algebraic equations (DAEs)

Often (not always) *sparse*.



Matrix




Graph


Consider ODEs  $\dot{x} = f(x)$  (special case  $f(x) = Ax$ ).

- Dependency graph: edge  $(i, j)$  if  $f_j$  depends on  $x_i$
- Sparsity means each  $f_j$  depends on only a few  $x_i$
- Often arises from physical or logical locality
- Corresponds to  $A$  being sparse (mostly zeros)

# Sparsity and Partitioning



Matrix



Graph

Want to partition sparse graphs so that

- Subgraphs are same size (load balance)
- Cut size is minimal (minimize communication)

We'll talk more about this later.

Consider ODEs  $\dot{x} = f(x)$  (special case  $f(x) = Ax$ ).

Might want  $f(x_*) = 0$ .

- Boils down to  $Ax = b$  (e.g. for Newton-like steps)
- Can solve directly or iteratively
- Sparsity matters a lot!

Consider ODEs  $\dot{x} = f(x)$  (special case  $f(x) = Ax$ ).

Might want  $x(t)$  for many  $t$  given  $x_0$

- Involves time-stepping (explicit or implicit)
- Implicit methods involve linear/nonlinear solves
- Need to understand stiffness and stability issues



Consider ODEs  $\dot{x} = f(x)$  (special case  $f(x) = Ax$ ).

Might want eigenvalues/vectors of  $A$  or  $f'(x_*)$ .

## Explicit Time Stepping

- Example: forward Euler:  $x_{k+1} = x_k + (\Delta t)f(x_k)$
- Next step depends only on earlier steps
- Simple algorithms
- May have stability issues with stiff systems

# Implicit Time Stepping

- Example: backward Euler:  $x_{k+1} = x_k + (\Delta t)f(x_{k+1})$
- Next step depends on itself and on earlier steps
- Algorithms involve solves — complication, communication!
- Larger time steps, each step costs more

In all cases, lots of time in sparse matvec:

- Iterative linear solvers: repeated sparse matvec
- Iterative eigensolvers: repeated sparse matvec
- Explicit time marching: matvecs at each step
- Implicit time marching: iterative solves (involving matvecs)

We need to figure out how to make matvec fast!

- Sparse matrix  $\implies$  mostly zero entries
  - Can also have “data sparseness” — representation with less than  $O(n^2)$  storage, even if most entries nonzero
- Could be implicit (e.g. directional differencing)
- Sometimes explicit representation is useful
- Easy to get lots of indirect indexing!
- Compressed sparse storage schemes help

## Example: Compressed Sparse Row




Figure 1: Illustration of compressed sparse row format

This can be even more compact:

- Could organize by blocks (block CSR)
- Could compress column index data (16-bit vs 64-bit)
- Various other optimizations — see OSKI

- ODE and DAE models widely used in engineering
- Different analyses: static, dynamic, modal
- Sparse linear algebra is often key

## Distributed Parameter Models

---



## Types of PDEs

Type	Example	Time?	Space dependence?
Elliptic	electrostatics	steady	global
Hyperbolic	sound waves	yes	local
Parabolic	diffusion	yes	global


Different types involve different communication:

- Global dependence  $\implies$  lots of communication (or tiny steps)
- Local dependence from finite wave speeds; limits communication

## Example: 1D Heat Equation

Consider flow (e.g. of heat) in a uniform rod

- Heat ( $Q$ )  $\propto$  temperature ( $u$ )  $\times$  mass ( $\rho$ )
- Heat flow  $\propto$  temperature gradient (Fourier's law)



## Example: 1D Heat Equation

Consider flow (e.g. of heat) in a uniform rod

- Heat ( $Q$ )  $\propto$  temperature ( $u$ )  $\times$  mass ( $\rho$ )
- Heat flow  $\propto$  negative temperature gradient (Fourier's law)

$$\begin{aligned}\frac{\partial Q}{\partial t} &\propto h \frac{\partial u}{\partial t} \\ &\approx C \left[ \frac{u(x-h) - u(x)}{h} + \frac{u(x+h) - u(x)}{h} \right] \\ &= C \left[ \frac{u(x-h) - 2u(x) + u(x+h)}{h^2} \right] \rightarrow C \frac{\partial^2 u}{\partial x^2}\end{aligned}$$

Heat equation with  $u(0) = u(1) = 0$ .

$$\frac{\partial u}{\partial t} = C \frac{\partial^2 u}{\partial x^2}$$

Spatial semi-discretization (second-order finite difference):

$$\frac{\partial^2 u}{\partial x^2} \approx \frac{u(x-h) - 2u(x) + u(x+h)}{h^2}$$

Yields system of ODEs (“method of lines”):

$$\frac{du}{dt} = -Ch^{-2}Tu$$
$$T = \begin{bmatrix} 2 & -1 & & & \\ -1 & 2 & -1 & & \\ & \ddots & \ddots & \ddots & \\ & & -1 & 2 & -1 \\ & & & -1 & 2 \end{bmatrix}$$

Now need to time step!


- Simplest scheme is Euler:

$$u(t + \Delta t) \approx u(t) + u'(t)\Delta t = (I - Ch^2T)u(t)$$

- Time step  $\equiv$  sparse matvec with  $(I - Ch^2T)$
- This may not end well...




## Explicit Data Dependence




Nearest neighbor interactions per step  $\implies$   
finite rate of numerical information propagation

## Explicit Time Stepping in Parallel




```
for t = 1 to N  
  communicate boundary data ("ghost cell")  
  take time steps locally  
end
```

# Overlapping Communication with Computation




```
for t = 1 to N
  start boundary data sendrecv
  compute new interior values
  finish sendrecv
  compute new boundary values
end
```

## Batching Time Steps



```
for t = 1 to N by B
  start boundary data sendrecv (B values)
  compute new interior values
  finish sendrecv (B values)
  compute new boundary values
end
```

# Explicit Pain



- Unstable for  $\Delta t > O(h^2)$
- Generally happens for parabolic (diffusive) equations
- But these ideas are great for hyperbolic equations!

- Backward Euler:  $u(t + \Delta t) \approx u(t) + \dot{u}(t + \Delta t)$
- Discretized time step:  $u(t + \Delta T) = (I + Ch^2T)^{-1}u(t)$
- No time step restriction for stability (good!)
- But each step involves a linear solve (not so good!)
  - Good if you like numerical linear algebra?

Explicit:

- Propagates information at finite rate
- Steps look like sparse matvec (in linear case)
- Stable step determined by fastest time scale
- Works fine for *hyperbolic* PDEs



Implicit:

- No need to resolve fastest time scales
- Steps can be long... but expensive
  - Linear/nonlinear solves at each step
  - Often these solves involve sparse matvecs
- Critical for parabolic PDEs

Consider 2D Poisson

$$-\nabla^2 u = -\frac{\partial^2 u}{\partial x^2} - \frac{\partial^2 u}{\partial y^2} = f$$

- Prototypical elliptic problem (steady state)
- Similar to a backward Euler step on heat equation

$$u_{i,j} = h^{-2} (4u_{ij} - u_{i-1,j} - u_{i+1,j} - u_{i,j-1} - u_{i,j+1})$$

## Second-Order Finite Differences

$$L = \left[ \begin{array}{ccc|cc|ccc} 4 & -1 & & -1 & & & & \\ -1 & 4 & -1 & & -1 & & & \\ & -1 & 4 & & & -1 & & \\ \hline -1 & & & 4 & -1 & & -1 & \\ & -1 & & -1 & 4 & -1 & & -1 \\ & & -1 & & -1 & 4 & & \\ \hline & & & -1 & & & 4 & -1 \\ & & & & -1 & & -1 & 4 & -1 \\ & & & & & -1 & & -1 & 4 \end{array} \right]$$

$N = n^d$  total unknowns

Ref: Demmel, *Applied Numerical Linear Algebra*, SIAM, 1997.

Method	Time	Space
Dense LU	$N^3$	$N^2$
Band LU	$N^2$	$N^{3/2}$
Jacobi	$N^2$	$N$
Explicit inv	$N^2$	$N^2$

Method	Time	Space
CG	$N^{3/2}$	$N$
Red-black SOR	$N^{3/2}$	$N$
Sparse LU	$N^{3/2}$	$N \log N$
FFT	$N \log N$	$N$
Multigrid	$N$	$N$

- Implicit solves or steady state  $\implies$  solving systems
- Nonlinear solvers generally linearize
- Linear solvers can be
  - Direct (hard to scale)
  - Iterative (often problem-specific)
- Iterative solves boil down to matvec!



Can be implicit or explicit (as with ODEs)

- Explicit (sparse matvec) — fast, but short steps?
  - works fine for hyperbolic PDEs
- Implicit (sparse solve)
  - Direct solvers are hard!
  - Sparse solvers turn into matvec again

Differential operators turn into local mesh stencils

- Matrix connectivity looks like mesh connectivity
- Can partition into subdomains that communicate only through boundary data
- More on graph partitioning later

Not all nearest neighbor ops are equally efficient!

- Depends on mesh structure
- Also depends on flops/point

- Next week: Distributed memory with MPI
- HW1 is posted: please run on Perlmutter!